

AI for Humanity Compass

Practical Guide for Responsible AI Decisions

Itamar Shabtai, Ph.D.
October 2025

“Management is doing things right; leadership is doing the right things.”

Peter Drucker

Table of Contents

1. What is the AI for Humanity Compass?
2. The Problem the Compass Solves
3. The Eight Compass Dimensions
4. How to Use the Compass - *Step-by-Step Guide*
5. What Good Looks Like - *Example: AI Triage Agent*
6. Why the Compass Matters
7. How to Use the Compass in Your Organization

1. What is the AI for Humanity Compass?

The AI for Humanity Compass is a **decision tool for AI projects**. Its goal is simple: keep every AI initiative aligned with **human dignity, safety, and well-being** - before launch, during deployment, and after it is live.

It exists because AI moves fast, but people, organizations, and policy move slow. The Compass is designed to close that gap: it helps teams slow down at the right moments, ask the right questions, and take responsibility for impact - not just capability, efficiency, or cost.

The Compass gives teams a shared language to answer eight essential questions:

1. Are we solving a real human need?
2. Who benefits and who could be harmed?
3. Are we living our values in the design?
4. Where can things go wrong, and how are we preventing that?
5. Where does a real human stay in control?
6. Are we respecting privacy and data boundaries?
7. What good are we trying to create for people, the organization, and society?
8. How will we measure trust, fairness, and well-being over time?

It is not theory. It is meant to be a **practical tool**.

2. The Problem the Compass Solves

The Compass exists to bridge a growing gap: AI technology is accelerating faster than our institutions, policies, cultures, and people can adapt. Teams can now build and ship powerful systems in days, but society - workers, patients, citizens, regulators - moves much more slowly in its ability to absorb consequences. The result is pressure to deploy before we know if something is safe, fair, or aligned with human dignity.

Most AI projects optimize for capability (“What can we build?”) and efficiency (“How much can we save?”).

What gets missed:

1. **Human impact is unclear.** Teams cannot clearly say who is helped, who is stressed, and who is exposed to harm.
2. **Values are not operational.** “We care about fairness” appears in slides, but there is no test, no owner, and no review cycle.
3. **No safety gates.** Problems are discovered only after a complaint, a headline, or an injury.
4. **No dignity check.** Stress, pressure, surveillance, bias, and exclusion are not treated as measurable outcomes.
5. **No escalation path.** The system produces confident answers, but it’s unclear when a real human takes responsibility.
6. **Unprepared humans in the loop.** Even when “a human will monitor,” that person often doesn’t have the authority, training, or capacity to intervene.

The Compass exists to surface these gaps early, before deployment, so leaders can act on them while they can still be fixed.

3. The Eight Compass Dimensions

1. **N – Purpose:** The human need. Why this matters to a person, not just to the business.
2. **NE – People:** Who we help, who we might hurt, and who we are responsible for protecting.
3. **E – Values:** The human values we refuse to trade away, and how we build them in.
4. **SE – Risks:** The main ways this can go wrong, and what we are doing about it.
5. **S – Human-in-the-Loop:** Where a real human can play a role, and Do we have the right people (skills, authority, capacity, clinical/operational judgment) in that loop.
6. **SW – Data & Privacy:** Data boundaries, consent, retention, and access.
7. **W – Outcomes:** The positive impact we are promising to deliver.
8. **NW – Metrics of Humanity:** How we will track trust, dignity, fairness, and well-being.




Figure 1: AI for Humanity Compass

4. How to Use the Compass - *Step-by-Step Guide*

Use the Compass any time you define, approve, pilot, or scale an AI system.

Step 1: Write the goal

Fill in one sentence: “We are building _____ so that _____ can _____.”

This locks the work to a human need, not just an efficiency target.

Step 2: Walk the Compass

Answer each point with short bullets:

1. **Purpose (N):** What real human need are we serving? What is out of scope?
2. **People (NE):** Who benefits? Who could be harmed, excluded, or stressed?
3. **Values (E):** Which values matter most (for example: safety, fairness, transparency, dignity)? How are we enforcing each one in practice?
4. **Risks (SE):** Top 2–3 things that can go wrong (misuse, bias, over-automation, unsafe advice). What is our control for each?
5. **Human-in-the-Loop (S):** Where must a qualified human review, approve, or override the AI? Do we have the right people (skills, authority, capacity, clinical/operational judgment) in that loop, and what is the escalation path?
6. **Data & Privacy (SW):** What is the minimum data we will collect? How do we get consent, limit retention, and control access?
7. **Outcomes (W):** What good should this create for (a) people, (b) the organization, and (c) society / access / equity?
8. **Metrics of Humanity (NW):** How will we measure trust, clarity, fairness, accessibility, and well-being over time?

Step 3: Traffic-light it

For each of the eight points, mark:

- Green = clear and acceptable
- Yellow = needs work
- Red = unsafe / unclear / not acceptable

Rule: If any point is Red, you do not ship.

Step 4: Decide “Go/Fix/Pause”

- **Go** = All items Green or Yellow, and the Yellow items have owners and deadlines.
- **Fix** = Too many Yellows, or missing owners.
- **Pause** = We identified a dignity/safety risk we are not willing to accept.

Record owners, deadlines, review dates, and escalation responsibilities.

5. What Good Looks Like - *Example: AI Triage Agent*

Project: Multilingual AI triage assistant for a clinic that helps patients describe their symptoms and get routed to the right level of care.

1. **Purpose:** Reduce confusion and anxiety for patients seeking care and help nurses focus on true urgent cases. Out of scope: diagnosis or treatment advice.
2. **People:** Helps patients (especially stressed, non-native speakers); supports nurses. At-risk: older adults, low-literacy patients. Duty: never leave high-risk patients alone with only AI.
3. **Values:**
 - a. Safety → nurse escalation on red flags
 - b. Empathy → calm, respectful language
 - c. Transparency → “You are chatting with an AI assistant, not a clinician.”
4. **Risks:**
 - a. Wrong or unsafe advice → restricted scripts + mandatory nurse review
 - b. Language bias → bilingual testing and parity tracking
 - c. Over-automation → human review required for high-risk cases
5. **Human-in-the-Loop:** Any red flag or uncertainty is sent to a licensed nurse within minutes. The nurse can override routing. We confirm that the nurses in the loop have the right training, authority, capacity, and emotional readiness to make those calls. They are not just “button clickers,” they are actual decision-makers.

6. **Data & Privacy:** Only collects symptoms + contact preference. Access is role-based. Chat logs auto-delete after 30 days.

7. **Outcomes:**

- a. Better care navigation: patients feel calmer and know next steps
- b. Reduced burden: nurses get fewer repeat “what do I do now?” calls
- c. Faster routing of urgent cases

8. **Metrics of Humanity:**

- a. Trust/clarity: % of patients who say “I understand what to do next”
- b. Equity: time-to-escalation gap between English and non-English speakers
- c. Accessibility: interface passes basic readability and accessibility checks

Step 3. Traffic-light (for this example)

- Purpose: ●
- People (vulnerable groups identified and duty of care named): ●
- Values (safety, empathy, transparency are enforced, not slogans): ●
- Risks (mitigations exist with owners): ● language parity still being tested
- Human-in-the-Loop (qualified nurses, with authority and override power): ●
- Data & Privacy (minimal PHI, retention limits, role-based access): ●
- Outcomes (clear patient and nurse benefit): ●
- Metrics of Humanity (defined, but some baselines not yet collected): ●

Rule reminder: Any Red would block deployment.

Step 4. Decide “Go/Fix/Pause”

Decision: FIX → GO

- We are not Red anywhere, but we have two Yellows:
 1. We still need baseline escalation parity between English and non-English speakers.
 2. We still need baseline scores for “I understand what to do next.”
- Owners are assigned (ML lead for language parity; Patient Experience lead for clarity scoring) and deadlines are set.

After those two fixes: status moves to **GO**.

This is what a responsible “Go” looks like: clear purpose, named protections, and ongoing measurement.

6. Why the Compass Matters

The Compass matters because AI today is deployed faster than people, policies, and safeguards can adapt. Most AI projects are pushed forward on speed and efficiency, while questions like “Who could be harmed?” and “Who is accountable if this goes wrong?” are left vague or ignored. The Compass forces those questions into the open before launch. It makes you define the real human benefit, identify who is at risk, prove that a qualified human can step in and stop harm, and commit to measuring trust, fairness, dignity, and well-being, not just performance and cost. It also gives leadership a clear, defensible moment to say “Pause” when the risk to people is unacceptable.

In practice, the Compass:

- Stops “ship now, clean up later.”
- Protects vulnerable groups, not just the average user.
- Turns ethics into ownership and action.
- Gives leaders permission to pause before harm, not after.

7. How to Use the Compass in Your Organization

The AI for Humanity Compass is not just a checklist — it’s how you control risk, protect dignity, and earn trust before you deploy AI that touches real people. You should run the Compass for any AI system that affects access, eligibility, safety, health, money, surveillance, or that replaces human judgment. It must be completed by a real cross-functional group (product, AI/tech lead, domain expert such as a nurse or HR lead, legal/privacy, UX/accessibility, and the actual human who will be “in the loop”), not by one person alone. You run it at three critical gates — before you build (“Should we even do this?”), before pilot (“Is this safe to try on real people/data?”),

and before scale (“Are we staffed and accountable to run this for real?”). The output becomes part of the official project record, just like a PRD or model card, and you revisit it whenever you expand the use case, touch a new population, or reduce/automate human review. You are doing this right when teams can clearly say who benefits, who is at risk, and how those at-risk groups are being protected; when “human in the loop” is a trained, empowered role and not a slide; and when leadership is willing to pause an AI system until dignity, safety, or equity concerns are fixed.

Your call to action:

- Treat the Compass as a launch gate.
- Put the *right* humans in the loop - with authority, skills, and capacity.
- Track humanity metrics (trust, clarity, equity, well-being, accessibility) next to cost and performance, not after.
- Be willing to say **Pause** when the cost to human dignity is too high.

***“The greatest danger in times of turbulence is not the turbulence,
it is to act with yesterday’s logic”, Peter Drucker***